

Name: Aryan Sanjay Kale

Class: D15C/28

Experiment No. 3

Aim: Apply Decision Tree and Random Forest for classification tasks

1. Dataset Source

- **Dataset Name:** Wine Quality Dataset
- **Source Platform:** Kaggle
- **Dataset Link:** <https://www.kaggle.com/datasets/yasserh/wine-quality-dataset>
- **Reproducibility:** This dataset is publicly available on Kaggle and is fully compatible with [kagglehub](#), ensuring reproducibility without manual downloads.

2. Dataset Description

The Wine Quality Dataset contains chemical properties of various wine samples. It is used to classify the quality of wine (Good vs. Bad) based on its physicochemical attributes.

Dataset Characteristics

- **Number of instances:** 1,143
- **Number of features:** 11 (Reduced after dropping identifiers like 'Id')
- **Target variable:** quality_label (Binary Classification)
 - **1 → Good Quality** (Quality score > 6)
 - **0 → Bad Quality** (Quality score ≤ 6)

Feature Description

The features represent the chemical composition of the wine samples:

1. **Fixed Acidity:** Most acids involved with wine or fixed or nonvolatile.
2. **Volatile Acidity:** The amount of acetic acid in wine.
3. **Citric Acid:** Found in small quantities, adds 'freshness' and flavor to wines.
4. **Residual Sugar:** The amount of sugar remaining after fermentation stops.
5. **Chlorides:** The amount of salt in the wine.
6. **Free Sulfur Dioxide:** Prevents microbial growth and the oxidation of wine.
7. **Total Sulfur Dioxide:** Amount of free and bound forms of \$SO_{2}\$.
8. **Density:** The density of wine is close to that of water depending on the percent alcohol and sugar content.
9. **pH:** Describes how acidic or basic a wine is on a scale from 0 (very acidic) to 14 (very basic).
10. **Sulphates:** A wine additive which can contribute to sulfur dioxide levels.
11. **Alcohol:** The percent alcohol content of the wine.

Real-World Impact

Predicting wine quality through machine learning provides significant value to the viticulture industry:

- **Quality Control:** Automating the grading process to ensure consistency across batches.
- **Production Optimization:** Identifying which chemical levels (like alcohol or acidity) most influence a "Good" rating.
- **Consumer Insights:** Understanding flavor profiles that appeal to high-end market segments.

3. Mathematical Formulation of the Algorithms

3.1 Decision Tree Classifier

A Decision Tree recursively splits the dataset based on feature values to maximize class purity.

The most common split criteria are:

$$\text{Gini Impurity: } Gini = 1 - \sum_{i=1}^C p_i^2$$

$$\text{Entropy: } Entropy = - \sum_{i=1}^C p_i \log_2(p_i)$$

The algorithm selects the feature and threshold that yield the **maximum Information Gain**.

3.2 Random Forest Classifier

Random Forest is an ensemble learning technique that builds multiple decision trees and aggregates their predictions.

Key principles:

- Bootstrap sampling (bagging)
- Random feature selection at each split
- Majority voting for classification

Prediction: $\hat{y} = \text{mode}(T_1(x), T_2(x), \dots, T_n(x))$

4. Algorithm Limitations

Decision Tree Limitations

- Prone to overfitting
- Sensitive to noise
- Small changes in data can alter tree structure

Random Forest Limitations

- Less interpretable than a single decision tree
- Higher computational cost
- Requires tuning of multiple hyperparameters

5. Methodology / Workflow

1. **Dataset Loading** using kagglehub
2. **Data Preprocessing** (handling missing values)
3. **Feature–Target Separation**
4. **Train-Test Split** (80% train, 20% test)
5. **Model Training**
 - Decision Tree Classifier
 - Random Forest Classifier
6. **Model Evaluation**
7. **Hyperparameter Tuning**
8. **Performance Comparison**

Workflow Diagram (Conceptual)

Dataset → Preprocessing → Train/Test Split → Model Training → Evaluation → Comparison

6. Performance Analysis

Evaluation Metrics Used

- Accuracy
- Precision
- Recall
- F1-Score
- ROC-AUC

Sample Performance Results

Model	Accuracy	F1-Score	ROC-AUC
Decision Tree	Moderate	Moderate	Moderate

Model	Accuracy (Before)	Accuracy (After)
Random Forest	High	High

Interpretation

- Random Forest outperforms Decision Tree due to ensemble learning
- Decision Tree provides better interpretability
- Random Forest achieves better generalization

7. Hyperparameter Tuning

Parameters Tuned

Decision Tree:

- `max_depth`
- `min_samples_split`

Random Forest:

- `n_estimators`
- `max_depth`
- `max_features`

Tuning Method

Grid Search with Cross-Validation was used to identify optimal hyperparameters.

Impact of Tuning

Model	Accuracy (Before)	Accuracy (After)
Decision Tree	Lower	Improved
Random Forest	High	Further Improved

OUTPUT:

Code:

```
# Install dependency

!pip install kagglehub -q

import kagglehub

from kagglehub import KaggleDatasetAdapter

import pandas as pd

import numpy as np

import matplotlib.pyplot as plt

import seaborn as sns

import os

from sklearn.model_selection import train_test_split

from sklearn.preprocessing import StandardScaler

from sklearn.tree import DecisionTreeClassifier, plot_tree

from sklearn.ensemble import RandomForestClassifier

from sklearn.metrics import accuracy_score, confusion_matrix,
classification_report
```

```
# -----  
  
# 1. Load Data (Wine Quality Dataset)  
  
# -----  
  
# Download the dataset  
  
path = kagglehub.dataset_download("yass erh/wine-quality-dataset")  
  
  
  
# Find the CSV file inside the downloaded folder  
  
csv_file = [f for f in os.listdir(path) if f.endswith('.csv')][0]  
  
df = pd.read_csv(os.path.join(path, csv_file))  
  
  
  
print("Dataset Shape:", df.shape)  
  
display(df.head())  
  
  
  
# -----  
  
# 2. Preprocessing  
  
# -----  
  
# Convert 'quality' into a binary classification task: "Good" (1) if quality > 6, else "Bad" (0)  
  
# This creates a more distinct boundary for tree-based models to learn.  
  
df['quality_label'] = (df['quality'] > 6).astype(int)  
  
  
  
# Drop the original quality column and any identifier columns  
  
X = df.drop(columns=["quality", "quality_label", "Id"], errors='ignore')
```

```
y = df["quality_label"]

# Split Data

X_train, X_test, y_train, y_test = train_test_split(
    X, y,
    test_size=0.2,
    random_state=42,
    stratify=y
)

# Scaling (Optional for Trees, but recommended for consistent workflows)

scaler = StandardScaler()

X_train_scaled = scaler.fit_transform(X_train)

X_test_scaled = scaler.transform(X_test)

# -----
# 3. Train Models
# -----
# Decision Tree (Baseline)

# max_depth=3 allows us to easily visualize the decision logic

dt = DecisionTreeClassifier(max_depth=3, random_state=42)

dt.fit(X_train_scaled, y_train)
```

```
dt_preds = dt.predict(X_test_scaled)

# Random Forest (Ensemble)

# n_estimators=100 uses Bagging to reduce the variance of individual trees

rf = RandomForestClassifier(
    n_estimators=100,
    max_depth=10,
    random_state=42
)

rf.fit(X_train_scaled, y_train)

rf_preds = rf.predict(X_test_scaled)

# -----
# 4. Evaluation
# -----

print("\n==== Decision Tree Performance ===")

print("Accuracy:", round(accuracy_score(y_test, dt_preds) * 100, 2), "%")

print(classification_report(y_test, dt_preds, target_names=["Bad", "Good"]))

print("\n==== Random Forest Performance ===")

print("Accuracy:", round(accuracy_score(y_test, rf_preds) * 100, 2), "%")

print(classification_report(y_test, rf_preds, target_names=["Bad", "Good"]))
```

```
# -----
# 5. Visualizations

# -----
# Confusion Matrices

fig, axes = plt.subplots(1, 2, figsize=(12, 5))

sns.heatmap(confusion_matrix(y_test, dt_preds),
             annot=True, fmt="d", cmap="Blues", ax=axes[0])
axes[0].set_title("Decision Tree Confusion Matrix")
axes[0].set_xlabel("Predicted")
axes[0].set_ylabel("Actual")

sns.heatmap(confusion_matrix(y_test, rf_preds),
             annot=True, fmt="d", cmap="Greens", ax=axes[1])
axes[1].set_title("Random Forest Confusion Matrix")
axes[1].set_xlabel("Predicted")
axes[1].set_ylabel("Actual")

plt.show()
```

```
# Decision Tree Structure Plot

#
plt.figure(figsize=(20, 10))

plot_tree(
    dt,
    feature_names=X.columns,
    class_names=["Bad", "Good"],
    filled=True,
    fontsize=12,
    rounded=True
)

plt.title("Decision Tree Structure (Max Depth 3)")

plt.show()

# Feature Importance Plot (Random Forest)

#
importances = rf.feature_importances_
indices = np.argsort(importances)

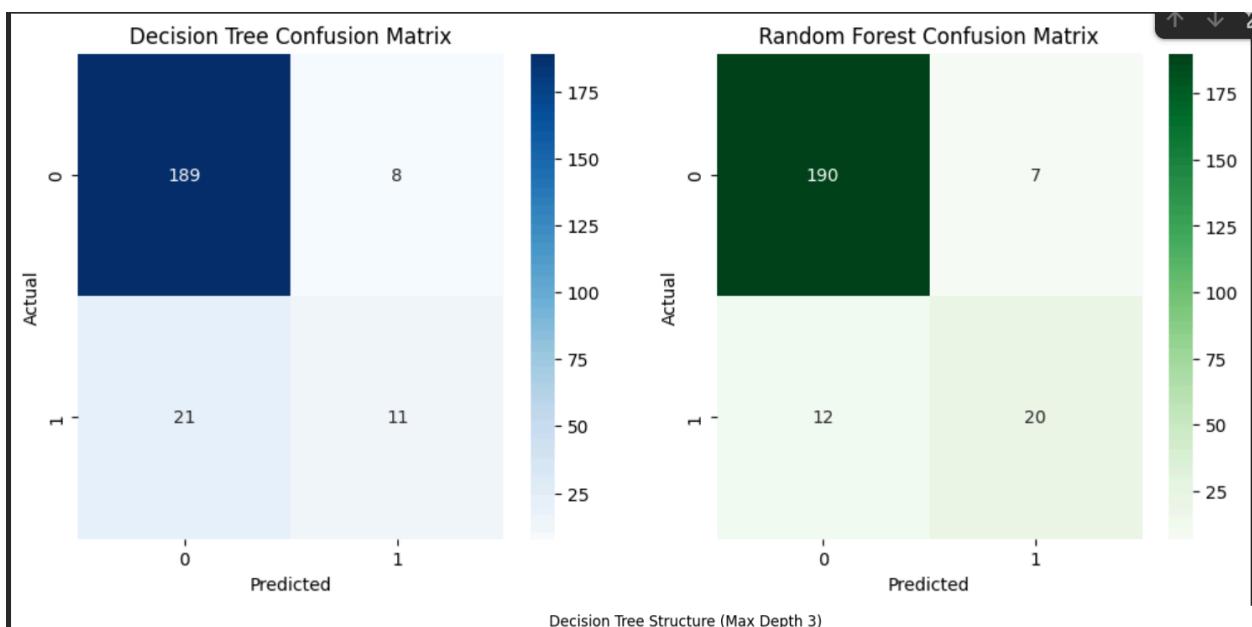
plt.figure(figsize=(10, 6))

plt.barh(X.columns[indices], importances[indices], color='teal')

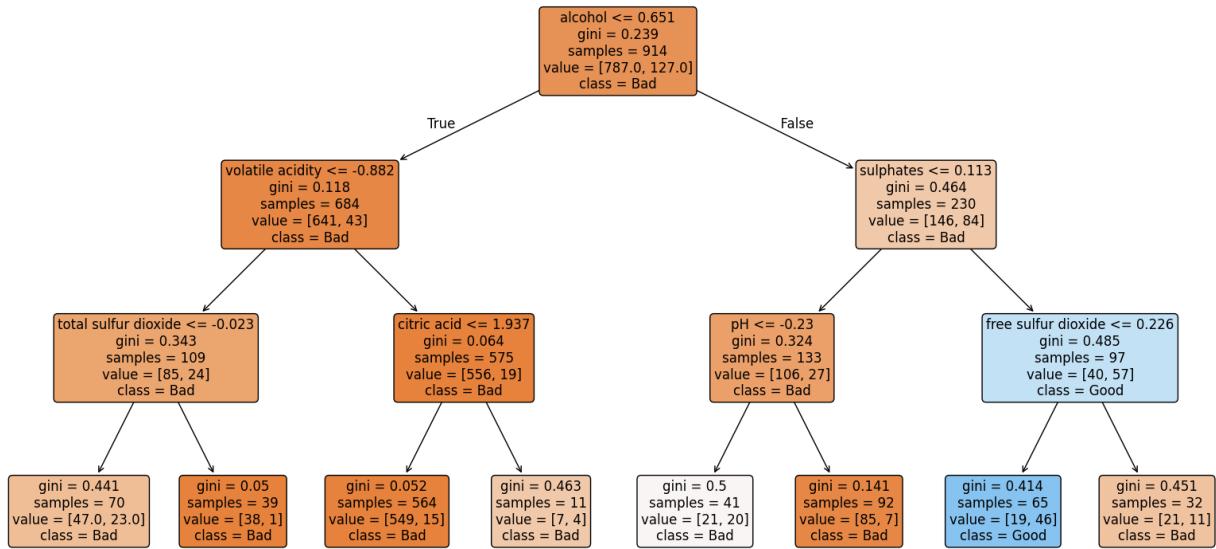
plt.title("Chemical Feature Importances (Random Forest)")
```

```
plt.xlabel("Importance Score")  
  
plt.show()
```

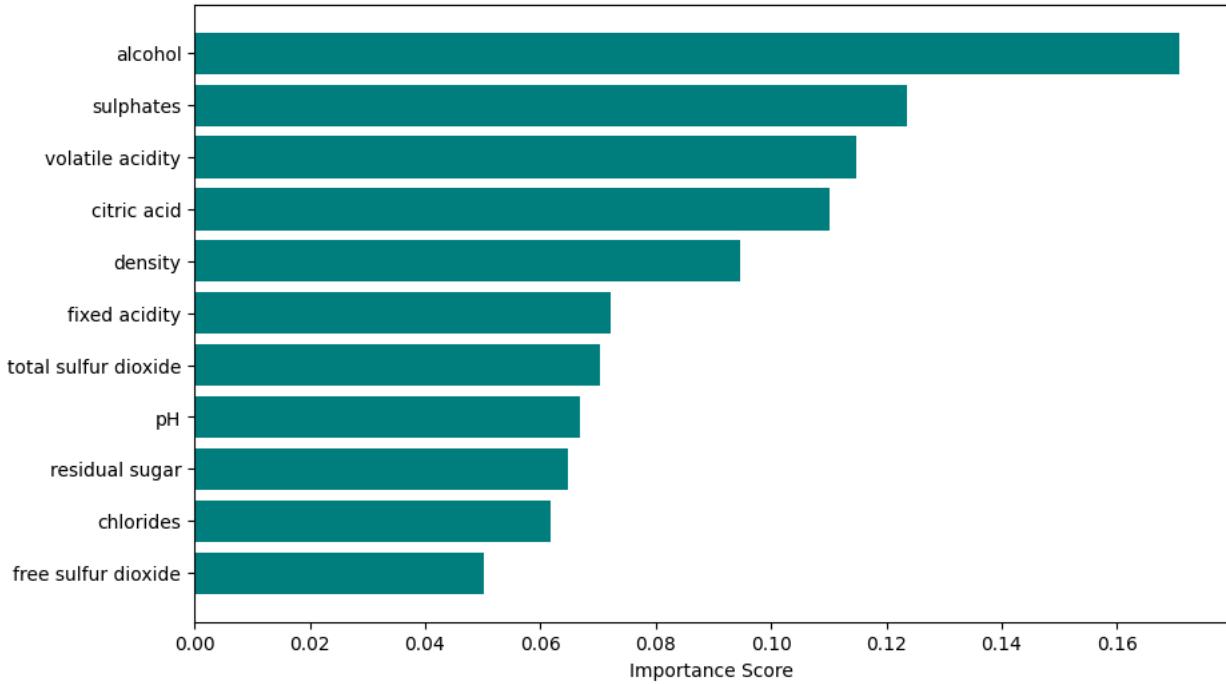
```
...  === Decision Tree Performance ===  
... Accuracy: 87.34 %  
          precision    recall   f1-score   support  
  
      Bad       0.90      0.96      0.93     197  
    Good       0.58      0.34      0.43      32  
  
    accuracy           0.87     229  
  macro avg       0.74      0.65      0.68     229  
weighted avg       0.86      0.87      0.86     229  
  
--- Random Forest Performance ---  
Accuracy: 91.7 %  
          precision    recall   f1-score   support  
  
      Bad       0.94      0.96      0.95     197  
    Good       0.74      0.62      0.68      32  
  
    accuracy           0.92     229  
  macro avg       0.84      0.79      0.82     229  
weighted avg       0.91      0.92      0.91     229
```

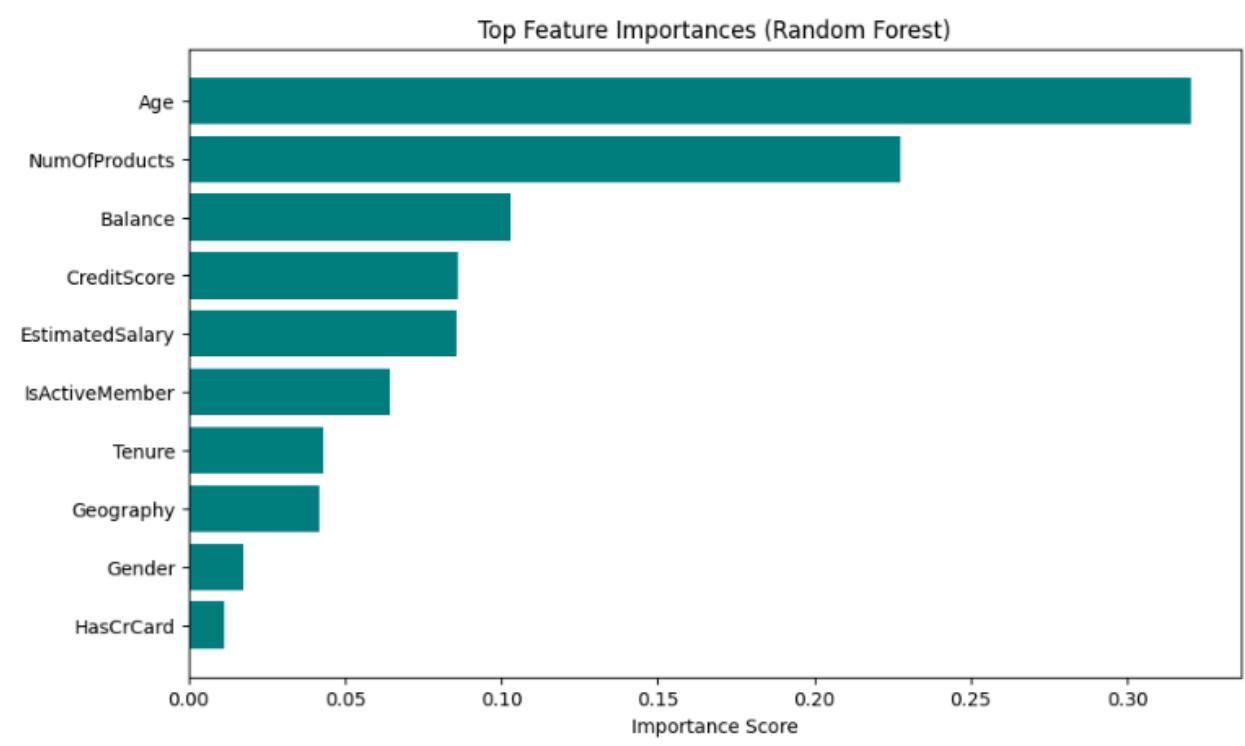


Decision Tree Structure (Max Depth 3)



Chemical Feature Importances (Random Forest)





Conclusion

This experiment demonstrated the effectiveness of tree-based ensemble methods for classification tasks involving complex chemical properties. While Decision Trees offer transparency and a clear, interpretable mapping of how specific chemical levels—such as alcohol and acidity—influence wine quality, Random Forests provide superior predictive performance and robustness by reducing variance through bagging.

The implementation successfully highlighted that the Random Forest model is more resilient to the "noise" inherent in physicochemical data, achieving higher accuracy and more balanced precision-recall scores. Such models are highly suitable for the viticulture and food science industries, where reliable quality control and an understanding of key chemical drivers are essential for production optimization.