

Name:**Aryan Sanjay Kale**

Class:**D15C** Roll no.**28**

Subject:**ML&DL**

Experiment No.4

AIM:Implement K-Nearest Neighbors (KNN) and evaluate model performance.

1. Dataset Source

- **Dataset Name:** Customer Personality Analysis
- **Source:** Kaggle - Customer Personality Analysis Dataset
- **Original Source:** Formatted for marketing research to provide a detailed analysis of a company's ideal customers.

2. Dataset Description

The dataset consists of various customer attributes including demographics, spending habits, and campaign responses. It is used to perform customer segmentation and predict spending behavior using distance-based classification.

- **Size:** 2,240 samples (rows) \times 29 columns (before preprocessing).
- **Target Variable:** Target (Binary Categorical)
 - 1 → **High Spender:** Total amount spent across all categories is greater than the median.
 - 0 → **Standard Spender:** Total amount spent is less than or equal to the median.
- **Key Features (Predictors):** The model utilizes numerical features to determine "neighbor" similarity:
 - Income:** The customer's annual household income (Critical for scaling).
 - Kidhome / Teenhome:** Number of children or teenagers in the customer's household.
 - Recency:** Number of days since the customer's last purchase.
 - MntWines / MntMeatProducts / etc.:** Amounts spent on different product categories.
 - Year_Birth:** Birth year of the customer to calculate age-based similarity.

3. Mathematical Formulation of the Algorithm

KNN is a non-parametric, lazy learning algorithm. It does not "learn" a model (like finding coefficients in regression); instead, it memorizes the training data.

A. Similarity Metric (Euclidean Distance)

To classify a new data point (x), the algorithm calculates its distance to every point in the training set ($x^{\{i\}}$). The most common metric is Euclidean Distance:

$$d(x, x^{(i)}) = \sqrt{\sum_{j=1}^n (x_j - x_j^{(i)})^2}$$

Where n is the number of features (4 in this case).

B. Classification Rule

1. Find the K nearest neighbors (points with the smallest distance d).
2. Assign the new point to the class that is most common (Mode) among those K neighbors.

$$\hat{y} = \text{mode}(y_1, y_2, \dots, y_K)$$

4. Algorithm Limitations

1. **Computational Cost:** It is "lazy," meaning all computation happens at prediction time. For large datasets, calculating the distance to *every* training point is slow.
2. **Sensitivity to Outliers:** If K is too small (e.g., $K=1$), a single mislabeled outlier can completely change the prediction.
3. **Scale Sensitivity:** KNN relies on distance. If one feature is measured in millimeters (e.g., 1000mm) and another in meters (e.g., 1m), the larger number will dominate the distance calculation. **Feature Scaling is mandatory.**

6. Code and Output

```

import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
import os
import kagglehub
from sklearn.model_selection import train_test_split
from sklearn.neighbors import KNeighborsClassifier
from sklearn.preprocessing import StandardScaler
from sklearn.metrics import accuracy_score, classification_report,
confusion_matrix

# =====
# 1. Load Data
# =====
# Download the Customer Personality Analysis dataset
path = kagglehub.dataset_download("imakash3011/customer-personality-analysis")
# The dataset is typically in a .tsv (tab-separated) format
csv_file = [f for f in os.listdir(path) if f.endswith('.csv') or
f.endswith('.tsv')][0]
df = pd.read_csv(os.path.join(path, csv_file), sep='\t')

# =====
# 2. Preprocessing
# =====
# Create a binary target: "High Spender" (1) if total spent > median, else (0)
spending_cols = ['MntWines', 'MntFruits', 'MntMeatProducts', 'MntFishProducts',
'MntSweetProducts', 'MntGoldProds']
df['Total_Spent'] = df[spending_cols].sum(axis=1)
df['Target'] = (df['Total_Spent'] > df['Total_Spent'].median()).astype(int)

# Select numerical features and drop rows with missing values (e.g., Income)
df_numeric = df.select_dtypes(include=[np.number]).dropna()

# Drop ID and target-related columns from features
X = df_numeric.drop(['ID', 'Total_Spent', 'Target'], axis=1)
y = df_numeric['Target']

# =====
# 3. Scaling (CRITICAL FOR KNN)

```

```

# =====
# KNN uses Euclidean distance; scaling prevents features with large ranges
# (like Income) from dominating features with small ranges (like Kidhome).
scaler = StandardScaler()
X_scaled = scaler.fit_transform(X)

# =====
# 4. Split Data
# =====
X_train, X_test, y_train, y_test = train_test_split(
    X_scaled, y,
    test_size=0.2,
    random_state=42,
    stratify=y
)

# =====
# 5. Train Baseline Model (K=5)
# =====
knn = KNeighborsClassifier(n_neighbors=5)
knn.fit(X_train, y_train)
y_pred = knn.predict(X_test)

# =====
# 6. Performance Metrics
# =====
print("--- Baseline KNN (K=5) Performance ---")
print(f"Accuracy: {accuracy_score(y_test, y_pred):.4f}")
print("\nClassification Report:\n", classification_report(y_test, y_pred,
target_names=['Standard Spender', 'High Spender']))

# Visualization: Confusion Matrix
plt.figure(figsize=(6, 5))
sns.heatmap(confusion_matrix(y_test, y_pred),
            annot=True,
            cmap='YlGnBu',
            fmt='d',
            xticklabels=['Standard', 'High'],
            yticklabels=['Standard', 'High'])
plt.title('Confusion Matrix: Customer Value Prediction (K=5)')
plt.ylabel('Actual Class')
plt.xlabel('Predicted Class')
plt.show()

# =====
# 7. Finding Optimal K (Elbow Method)
# =====
error_rate = []
for i in range(1, 20):
    knn_i = KNeighborsClassifier(n_neighbors=i)
    knn_i.fit(X_train, y_train)
    pred_i = knn_i.predict(X_test)
    error_rate.append(np.mean(pred_i != y_test))


plt.figure(figsize=(10, 6))

```

```

plt.plot(range(1, 20), error_rate, color='blue', linestyle='dashed', marker='o')
plt.title('Error Rate vs. K Value')
plt.xlabel('K')
plt.ylabel('Error Rate')
plt.show()

```



Typical Analysis:

- **Accuracy:** You will likely see an accuracy of **1.0 (100%)** or **0.96 (96.6%)**. This is because the Iris dataset features separate the classes very clearly.
- **Confusion Matrix:** Look for off-diagonal numbers. If there are any errors, they usually occur between *Versicolor* and *Virginica* because these two species look somewhat similar (their clusters overlap slightly), whereas *Setosa* is very distinct.

7. Hyperparameter Tuning (The Elbow Method)

Unlike regression where we tune alpha, in KNN we tune K (Number of Neighbors).

- Small K (e.g., 1): Low bias, High variance (Model is too jagged/sensitive).
- Large K: High bias, Low variance (Model is too smooth/simple).

We use the Elbow Method to find the sweet spot where the Error Rate is lowest.

```

# --- HYPERPARAMETER TUNING: ELBOW METHOD ---
error_rate = []

# Will take some time
for i in range(1, 20):

```


```

knn_i = KNeighborsClassifier(n_neighbors=i)
knn_i.fit(X_train, y_train)
pred_i = knn_i.predict(X_test)
# Calculate average error (mean of boolean array where pred != actual)
error_rate.append(np.mean(pred_i != y_test))

# Plot the Error Rate
plt.figure(figsize=(10, 6))
plt.plot(range(1, 20), error_rate, color='blue', linestyle='dashed',
marker='o',
markerfacecolor='red', markersize=10)
plt.title('Error Rate vs. K Value')
plt.xlabel('K')
plt.ylabel('Error Rate')
plt.show()

# Find minimum error
optimal_k = error_rate.index(min(error_rate)) + 1
print(f"Minimum Error found at K = {optimal_k}")

```



Interpretation:

- The Graph: You will see the line drop quickly.
- The Elbow: If the error is high at K=1 and drops at K=3, but stays flat after K=5, then K=3 or K=5 is the optimal choice (choosing the smaller, simpler number is usually better if performance is equal).

8. Conclusion

In this experiment, we implemented the K-Nearest Neighbors classifier on the Iris dataset. •

Performance: The model achieved an outstanding accuracy of [Insert Score, e.g., 100%], proving that the physical dimensions of Iris sepals and petals are highly predictive of their species. •

Importance of Scaling: Feature scaling was applied to ensure that petal length (which varies more) did not disproportionately influence the Euclidean distance calculation.

• Tuning: Using the Elbow Method, we determined that K=[Insert Optimal K] provided the most stable predictions, minimizing the risk of overfitting while maintaining maximum accuracy. KNN

proved to be a highly effective, albeit computationally intensive, algorithm for this classification task.